

S3-TASK

1. Create an S3 bucket and upload some objects to S3.

Login to Aws Console

Go to → S3 → Create Bucket → unique Name

uncheck the Block All Public

Create

upload Any file

Give ASL as Make Public

then Copy the Object link and paste it into the browser

The screenshot displays the AWS S3 console interface. The top navigation bar shows the user is logged in as 'Sunithmaz' in the 'United States (N. Virginia)' region. The main content area is divided into two sections: 'Upload' and 'Buckets'.

Upload Section:

- Upload Info:** A message states, 'Add the files and folders you want to upload to S3. To upload a file larger than 160GB, use the AWS CLI, AWS SDKs or Amazon S3 REST API. [Learn more](#)'.
- Files and folders (1 total, 5.1 MB):** A table lists the files to be uploaded. The first entry is 'windows-95-dark-3840x2160-27097.jpg' with a size of 5.1 MB and type 'image/jpeg'.
- Destination:** The destination is set to 's3://test-1342'.
- Destination details:** A note states, 'Bucket settings that impact new objects stored in the specified destination.'


Buckets Section:

- General purpose buckets (1):** A table lists the buckets. The first entry is 'test-1342' in the 'US East (N. Virginia) us-east-1' region, created on 'September 3, 2026, 14:45:29 (UTC+05:30)'.

Make public Section:

- Make public Info:** A message states, 'The make public action enables public read access in the object access control list (ACL) settings. [Learn more](#)'.
- Specified objects:** A table lists the objects to be made public. The first entry is 'windows-95-dark-3840x2160-27097.jpg' with a size of 5.1 MB and type 'jpg'.

S3-TASK



2. Deploy a static website in the S3 bucket.


Step 1 — Open Properties

Inside: → **S3** → **test-1342** → Click-**Properties**

Static website hosting → Enable website hosting →

Index document → index.html

For **Error document** → error.html



S3-TASK

aws [Search] [Alt+S] Ask Amazon Q United States (N. Virginia) Sumit

Amazon S3

Upload succeeded
For more information, see the [Files and folders](#) table.

Upload: status

After you navigate away from this page, the following information is no longer available.

Summary

Destination s3://test-1342	Succeeded 1 file, 188.0 B (100.00%)	Failed 0 files, 0 B (0%)
--------------------------------------	---	------------------------------------

Files and folders | Configuration

Files and folders (1 total, 188.0 B)

Find by name

Name	Folder	Type	Size	Status	Error
index.html.html	-	text/html	188.0 B	Succeeded	-

test-1342.s3.us-east-1.amazonaws.com/index.html.html

Movies MY Sites Ai's Gmail YouTube Maps Translate Resume Examples f

Hello from Amazon S3!

This website is hosted using Amazon S3.


Not secure test-1342.s3-website-us-east-1.amazonaws.com

Movies MY Sites Ai's Gmail YouTube Maps Translate Resume Examples f

under Maintaince!

This website is hosted using Amazon S3.

S3-TASK



3. Enable cross-region replication on S3 buckets.

In another Region

S3 → General purpose buckets → Create bucket → Unique name

Enable Versioning on source bucket

Create the replication rule

S3 → test-1342 → Management → Find-Replication rules → Create replication rule

Give the rule a name → Choose → Apply to all objects in the bucket

Choose destination → Choose a bucket in this account

IAM role → Create new role → Save

Test replication → Upload →

Check destination bucket → S3 → test-1342-replica → Objects

S3-TASK

Amazon S3 > Buckets > test-1342 > Replication rules > Create replication rule

Choose a rule scope

☐ Limit the scope of this rule using one or more filters

☒ Apply to all objects in the bucket

Choose a bucket

S3 Buckets

Buckets (1/2)

Name

AWS Region

☐ test-1342

US East (N. Virginia) us-east-1

☒ test-1342-replica

US West (Oregon) us-west-2

Cancel

Choose path

Replicate existing objects?

You can enable a one-time Batch Operations job from this replication configuration to replicate objects that already exist in the bucket and to synchronize the source and destination buckets. [Learn more](#) or [see pricing](#)

Existing objects

☒ No, do not replicate existing objects.

☐ Yes, replicate existing objects.

Cancel

Submit

Replication rules (1)

View details

Edit rule

Delete

Actions

Create replication rule

Use replication rules to define options you want Amazon S3 to apply during replication such as server-side encryption, replica ownership, transitioning replicas to another storage class, and more. Replication also includes any annotations attached to your objects. [Learn more](#)

Replication now includes replicating any annotations attached to your objects, which might increase your costs. To understand the potential cost impact, review your replication rules and annotations usage.

Replication rule name	Status	Destination bucket	Destination Region	Priority	Scope	Storage class	Replica owner	Replication Time Control	KMS-encrypted objects (SSE-KMS or DSSE-KMS)	Replica modification sync
☐ replicate-all-objects	☒ Enabled	s3://test-1342-replica	US West (Oregon) us-west-2	0	Entire bucket	Same as source	Same as source	Disabled	Do not replicate	Disabled

[View replication configuration](#)

S3-TASK

4. Configure a bucket policy so only the Admin user can see the objects of the S3 bucket.

Find the IAM user's ARN

IAM → Users → Devops → Permissions → add permissions → Attach policies directly
→ AmazonS3ReadOnlyAccess

Next → Add permissions

Block public access

S3 → Buckets → test-1342 → Permissions

Block public access

Edit → Enable: ☒ Block all public access

Replace the public bucket policy

test-1342 → Permissions → Edit

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Sid": "AllowAdminReadObjects",
      "Effect": "Allow",
      "Principal": {
        "AWS": "arn:aws:iam::123456789012:user/Admin"
      },
      "Action": "s3:GetObject",
      "Resource": "arn:aws:s3:::test-1342/*"
    },
    {
      "Sid": "AllowAdminListBucket",
      "Effect": "Allow",
      "Principal": {
        "AWS": "arn:aws:iam::123456789012:user/Admin"
      },
      "Action": "s3:ListBucket",
      "Resource": "arn:aws:s3:::test-1342"
    }
  ]
}
```

Sign in as your Admin IAM user.

I can see all the lists in the S3 bucket

Test with another user

another IAM user, sign in as that user.

The non-Admin user should not be able to read the objects.

You may see: Access Denied

S3-TASK

Add permissions

Add user to an existing group or create a new one. Using groups is a best-practice way to manage user's permissions by job functions. [Learn more](#)

Permissions options

☐ Add user to group

Add user to an existing group, or create a new group. We recommend using groups to manage user permissions by job function.

☐ Copy permissions

Copy all group memberships, attached managed policies, inline policies, and any existing permissions boundaries from an existing user.

☒ Attach policies directly

Attach a managed policy directly to a user. As a best practice, we recommend attaching policies to a group instead. Then, add the user to the appropriate group.

Permissions policies (1/1575)

Filter by Type

All types

1 match

☒ Policy name

Type

Attached entities

☒	AmazonS3ReadOnlyAccess	AWS managed	0
-------------------------------------	------------------------	-------------	---

Cancel Next

Amazon S3 > Buckets > test-1342 > Edit Block public access (bucket settings)

Edit Block public access (bucket settings)

Block public access (bucket settings)

Public access is granted to buckets and objects through access control lists (ACLs), bucket policies, access point policies, or all. In order to ensure that public access to all your S3 buckets and objects is blocked, you must turn on Block all public access. AWS recommends that you turn on Block all public access, but before applying any of these settings, ensure that your applications will work correctly without public access. If you require specific access points, you can turn off Block all public access and configure individual settings below to suit your specific storage use cases. [Learn more](#)

☒ Block all public access

Turning this setting on is the same as turning on all four settings below. Each of the following settings are independent of one another.

☒ Block public access to buckets and objects granted through new access control lists (ACLs)

S3 will block public access permissions applied to newly added buckets or objects, and prevent the creation of new public access ACLs for existing buckets and objects. This setting doesn't change existing public access permissions.

☒ Block public access to buckets and objects granted through any access control lists (ACLs)

S3 will ignore all ACLs that grant public access to buckets and objects.

☒ Block public access to buckets and objects granted through new public bucket or access point policies

S3 will block new bucket and access point policies that grant public access to buckets and objects. This setting doesn't change any existing policies that allow public access to S3 resources.

☒ Block public and cross-account access to buckets and objects through any public bucket or access point policies

S3 will ignore public and cross-account access for buckets or access points with policies that grant public access to buckets and objects.

aws [Search] [Alt+S]

Amazon S3 > Buckets > test-1342 > Edit bucket policy

Bucket policy

The bucket policy, written in JSON, provides access to the objects stored in the bucket. Bucket policies don't apply to objects owned by other accounts. [Learn more](#)

Bucket ARN

arn:aws:s3:::test-1342

Policy

```
1 - {
2   "Version": "2012-10-17",
3   "Statement": [
4     {
5       "Sid": "AllowAdminReadObjects",
6       "Effect": "Allow",
7       "Principal": {
8         "AWS": "arn:aws:iam::567487919764:user/DevopsA"
9       },
10      "Action": "s3:GetObject",
11      "Resource": "arn:aws:s3:::test-1342/*"
12    },
13    {
14      "Sid": "AllowAdminListBucket",
15      "Effect": "Allow",
16      "Principal": {
17        "AWS": "arn:aws:iam::567487919764:user/DevopsA"
18      },
19      "Action": "s3:ListBucket",
20      "Resource": "arn:aws:s3:::test-1342"
```

S3-TASK

The screenshot shows the Amazon S3 console interface. The left sidebar contains navigation links for Buckets, Files, Access management and security, and Storage management and security. The main content area displays the 'General purpose buckets (0)' section. A red error message is visible: 'You don't have permissions to list buckets. After you or your AWS administrator has updated your permissions to allow the s3:ListAllMyBuckets action, refresh this page. Learn more about Identity and access management in Amazon S3.' Below the error message, there are sections for 'Account snapshot' and 'External access summary'. The top of the console shows the AWS logo, search bar, and account information (Europe (Stockholm), Account ID: 5674-8791-9764).

Amazon S3 > Buckets

General purpose buckets (0) Info

Buckets are containers for data stored in S3.

Find buckets by name

Name | AWS Region | Creation date

You don't have permissions to list buckets

After you or your AWS administrator has updated your permissions to allow the s3:ListAllMyBuckets action, refresh this page. Learn more about [Identity and access management in Amazon S3](#)

Diagnose with Amazon Q

Account snapshot Info

Updated daily

View dashboard

Storage Lens provides visibility into storage usage and activity trends.

External access summary Info

Updated daily

External access findings help you identify bucket permissions that allow public access or access from other AWS accounts.

Europe (Stockholm) Account ID: 5674-8791-9764 Test

Account ID

5674-8791-9764

Account name

Access denied

Account color

Access denied

IAM user

Test

Account

Organization

Service Quotas

Billing and Cost Management

Security credentials

Console Mobile App

Agent Toolkit for AWS

Turn on multi-session support

S3-TASK

- Set up lifecycle policies to automatically transition or delete objects based on specific criteria.

AWS Console → S3 → Buckets → test-1342

Management → Find: Life cycle rules → Create lifecycle rule

Lifecycle rule name → test-1342-lifecycle

Choose a rule scope → Apply to all objects in the bucket

Configure transitions

Lifecycle rule actions:

☒ Transition current versions of objects between storage classes

First transition → S3 Standard-IA → Days after object creation 30 → After 30 days → S3 Standard-IA

Second transition → Glacier Flexible Retrieval → Days after object creation: 90

Configure deletion → Permanently delete noncurrent versions of objects

Lifecycle rule actions

☒ Permanently delete noncurrent versions of objects

Number of days after objects become noncurrent: 30

This means when an older version becomes noncurrent, S3 can permanently delete that version after 30 days.

Review the rule

Rule name:

test-1342-lifecycle

Scope:

All objects in bucket

Actions:

☒ Transition current versions

30 days → S3 Standard-IA

90 days → Glacier Flexible Retrieval

☒ Permanently delete noncurrent versions

30 days after becoming noncurrent

Amazon S3 > Buckets > test-1342

test-1342

Objects | Metadata | Properties | Permissions | Metrics | **Management** | File systems - new | Access Points

Lifecycle configuration

To manage your objects so that they are stored cost effectively throughout their lifecycle, configure their lifecycle. A lifecycle configuration is a set of rules that define actions that Amazon S3 applies to a group of objects. Lifecycle rules run once per day.

Lifecycle rules

Use lifecycle rules to define actions you want Amazon S3 to take during an object's lifetime such as transitioning objects to another storage class, archiving them, or deleting them after a specified period of time. [Learn more](#)

[View details](#) [Edit](#) [Delete](#) [Actions](#) [Create lifecycle rule](#)

Lifecycle rule name	Status	Scope	Current version actions	Noncurrent versions actions	Expired object delete mar...	Incomplete multipart upl...
No lifecycle rules There are no lifecycle rules for this bucket.						

[Create lifecycle rule](#)

S3-TASK

We don't recommend transitioning objects less than 128 KB because the transition costs can outweigh the storage savings. If your use case requires transitioning objects less than 128 KB, specify a minimum object size filter for each applicable lifecycle rule with a transition action.

Transition current versions of objects between storage classes

Choose transitions to move current versions of objects between storage classes based on your use case scenario and performance access requirements. These transitions start from when the objects are created and are consecutively applied. [Learn more](#)

Choose storage class transitions

Standard-IA ▼ Days after object creation 30 Remove

Glacier Instant Retrieval ▼ 90 Remove

Add transition

Review transition and expiration actions

Lifecycle rule actions

Choose the actions you want this rule to perform.

- ☒ Transition current versions of objects between storage classes
This action will move current versions.
- ☐ Transition noncurrent versions of objects between storage classes
This action will move noncurrent versions.
- ☐ Expire current versions of objects
- ☒ Permanently delete noncurrent versions of objects
- ☐ Delete expired object delete markers or incomplete multipart uploads
These actions are not supported when filtering by object tags or object size.

Transitions are charged per request
For a lifecycle transition action, each request corresponds to an object transition. For details on lifecycle transition pricing, see requests pricing info on the **Storage & requests** tab of the [Amazon S3 pricing page](#).

☒ I acknowledge that this lifecycle rule will incur a transition cost per request.

By default, objects less than 128KB will not transition across any storage class.

Lifecycle configuration

To manage your objects so that they are stored cost effectively throughout their lifecycle, configure their lifecycle. A lifecycle configuration is a set of rules that define actions that Amazon S3 applies to a group of objects. Lifecycle rules run once per day.

Default minimum object size for transitions
All storage classes 128K

Lifecycle rules (1)

Use lifecycle rules to define actions you want Amazon S3 to take during an object's lifetime such as transitioning objects to another storage class, archiving them, or deleting them after a specified period of time. [Learn more](#)

View details Edit Delete Actions ▼ Create lifecycle rule

Lifecycle rule name	Status	Scope	Current version actions	Noncurrent versions actions	Expired object delete mar...	Incomplete multipart upl...
☐ test-1342-lifecycle	Enabled	Entire bucket	Transition to Standard-IA, then i	Permanently delete	-	-

[View lifecycle configuration](#)

6. Push some objects to S3 using the AWS CLI.

aws configure

AWS Access Key ID [None]: AKIAYIIHKO2KH5IAAZOC

Secret access key : +akA/36e/KBfkyQ49GWXIORpsyzDlunyc

Default region name [None]:us-east-1

Default output format [None]:Json

Verify AWS identity

aws sts get-caller-identity

```
{
  "UserId": "567487919764",
  "Account": "567487919764",
  "Arn": "arn:aws:iam::567487919764:root"
}
```

Verify your bucket

aws s3 ls

S3-TASK

Create a local folder for the test

```
mkdir s3-cli-test → cd s3-cli-test
```

```
echo "Hello from AWS CLI" > hello.txt
```

```
echo ""This file was uploaded using AWS CLI " > cli-test.txt
```

```
echo "S3 Task 6 " > task6.txt
```

Upload ONE object

```
aws s3 cp hello.txt s3://test-1342/
```

You should see something similar to:

```
upload: ./hello.txt to s3://test-1342/hello.txt
```

```
aws s3 cp cli-test.txt s3://test-1342/
```

```
aws s3 cp task6.txt s3://test-1342/
```

Verify the objects

```
aws s3 ls s3://test-1342/
```

```
HP@Ashish MINGW64 ~
$ aws s3 ls
2026-09-03 14:45:32 test-1342

HP@Ashish MINGW64 ~
$ aws --version
aws-cli/2.36.32 Python/3.14.6 Windows/11 script-exe/AMD64

HP@Ashish MINGW64 ~
$ aws sts get-caller-identity
{
  "UserId": "567487919764",
  "Account": "567487919764",
  "Arn": "arn:aws:iam::567487919764:root"
}

HP@Ashish MINGW64 ~
$ BUCKET="test-1342"
```

```
HP@Ashish MINGW64 ~
$ BUCKET="test-1342"
```

```
HP@Ashish MINGW64 ~
$ REGION="us-east-1"
```

```
HP@Ashish MINGW64 ~
$ echo $BUCKET
test-1342
```

```
HP@Ashish MINGW64 ~
$ echo $REGION
us-east-1
```

```
HP@Ashish MINGW64 ~
$ aws s3api head-bucket --bucket "$BUCKET"
{
  "BucketArn": "arn:aws:s3:::test-1342",
  "BucketRegion": "us-east-1",
  "AccessPointAlias": false
}
```

```
HP@Ashish MINGW64 ~
$ mkdir s3-cli-test
```

```
HP@Ashish MINGW64 ~
$ cd s3-cli-test
```

```
HP@Ashish MINGW64 ~/s3-cli-test
$ pwd
/c/Users/HP/s3-cli-test
```

S3-TASK

```
HP@Ashish MINGW64 ~/s3-cli-test
$ echo "Hello from AWS CLI" > hello.txt

HP@Ashish MINGW64 ~/s3-cli-test
$ echo "This file was uploaded using AWS CLI" > hello.txt

HP@Ashish MINGW64 ~/s3-cli-test
$ ;s
bash: syntax error near unexpected token `;'

HP@Ashish MINGW64 ~/s3-cli-test
$ ls
hello.txt

HP@Ashish MINGW64 ~/s3-cli-test
$ echo "Hello from AWS CLI" >hello.txt

HP@Ashish MINGW64 ~/s3-cli-test
$ ls
hello.txt

HP@Ashish MINGW64 ~/s3-cli-test
$ cat hello.txt
Hello from AWS CLI

HP@Ashish MINGW64 ~/s3-cli-test
$ echo "This file was uploaded using AWS CLI" > cli-test.txt

HP@Ashish MINGW64 ~/s3-cli-test
$ echo "S3 Task 6" > task6.txt

HP@Ashish MINGW64 ~/s3-cli-test
$ ls
cli-test.txt  hello.txt  task6.txt
```

```
HP@Ashish MINGW64 ~/s3-cli-test
$ aws s3 cp hello.txt s3://test-1342/
upload: .\hello.txt to s3://test-1342/hello.txt


HP@Ashish MINGW64 ~/s3-cli-test
$ aws s3 cp cli-test.txt s3://test-1342/
upload: .\cli-test.txt to s3://test-1342/cli-test.txt

HP@Ashish MINGW64 ~/s3-cli-test
$ aws s3 cp cli-test.txt s3://test-1342/
upload: .\cli-test.txt to s3://test-1342/cli-test.txt

HP@Ashish MINGW64 ~/s3-cli-test
$ aws s3 ls s3://test-1342/
2026-09-03 15:56:03      162589 1aas.png
2026-09-03 17:30:18         37 cli-test.txt
2026-09-03 15:37:23     53261 error.html.html
2026-09-03 15:35:58     53261 error.html.txt
2026-09-03 17:29:46        19 hello.txt
2026-09-03 14:49:33   5378539 windows-95-dark-3840x2160-27097.jpg

HP@Ashish MINGW64 ~/s3-cli-test
$ !
```

S3-TASK



7. Write a Bash script to create an S3 bucket.

```
vim create-s3-bucket.sh
```

```
#!/bin/bash
```

```
# Bucket name
```

```
BUCKET_NAME="test-1342-script-2026"
```

```
# AWS region
```

```
REGION="us-east-1"
```

```
echo "Creating S3 bucket: $BUCKET_NAME"
```

```
echo "Region: $REGION"
```

```
aws s3api create-bucket \
  --bucket "$BUCKET_NAME" \
  --region "$REGION"
```

```
if [ $? -eq 0 ]; then
```

```
  echo "S3 bucket created successfully!"
```

```
Else
```

```
  echo "Failed to create S3 bucket."
```

```
  exit 1
```

```
Fi
```

```
echo "Verifying bucket..."
```

```
aws s3api head-bucket --bucket "$BUCKET_NAME"
```

```
if [ $? -eq 0 ]; then
```

```
  echo "Bucket verification successful."
```

```
Else
```

```
  echo "Bucket verification failed."
```

```
  exit 1
```

S3-TASK

fi

```
chmod +x create-s3-bucket.sh
```

```
ls -l create-s3-bucket.sh
```

```
./create-s3-bucket.sh
```

```
aws s3 ls
```

S3-TASK

```
#!/bin/bash

# Bucket name
BUCKET_NAME="test-1342-script-2026"

# AWS region
REGION="us-east-1"

echo "Creating S3 bucket: $BUCKET_NAME"
echo "Region: $REGION"

aws s3api create-bucket \
  --bucket "$BUCKET_NAME" \
  --region "$REGION"

if [ $? -eq 0 ]; then
  echo "S3 bucket created successfully!"
else
  echo "Failed to create S3 bucket."
  exit 1
fi

echo "Verifying bucket..."

aws s3api head-bucket --bucket "$BUCKET_NAME"

if [ $? -eq 0 ]; then
  echo "Bucket verification successful."
else
  echo "Bucket verification failed."
  exit 1
fi

~
```

```
HP@Ashish MINGW64 ~/s3-cli-test
$ ls -l create-s3-bucket.sh
-rwxr-xr-x 1 HP 197121 563 Sep  3 17:42 create-s3-bucket.sh*

HP@Ashish MINGW64 ~/s3-cli-test
$ ./create-s3-bucket.sh
Creating S3 bucket: test-1342-script-2026
Region: us-east-1
{
  "Location": "/test-1342-script-2026",
  "BucketArn": "arn:aws:s3:::test-1342-script-2026"
}

S3 bucket created successfully!
Verifying bucket...
{
  "BucketArn": "arn:aws:s3:::test-1342-script-2026",
  "BucketRegion": "us-east-1",
  "AccessPointAlias": false
}

Bucket verification successful.

HP@Ashish MINGW64 ~/s3-cli-test
$ aws s3 ls
2026-09-03 14:45:29 test-1342
2026-09-03 15:52:28 test-1342-replica
2026-09-03 17:42:56 test-1342-script-2026

HP@Ashish MINGW64 ~/s3-cli-test
$ |
```

S3-TASK

8. Upload a 1 GB file to S3 using the CLI.

`df -h` = Look at the available space for your main filesystem

`dd if=/dev/zero of=large-file-1GB.bin bs=1M count=1024` = to make 1 gb file

`ls -lh large-file-1GB.bin`

`du -h large-file-1GB.bin`

Make sure AWS CLI is working

`aws s3 ls s3://test-1342/`

Upload the 1 GB file

`aws s3 cp large-file-1GB.bin s3://test-1342/`

The command is:

`aws` → AWS CLI `s3` → S3 service `cp` → Copy/upload `large-file-1GB.bin` → Local 1 GB file `s3://test-1342/`
→ Your S3 bucket

Verify the upload

`aws s3 ls s3://test-1342/`

2026-09-03 17:30:18	37	cli-test.txt
2026-09-03 15:37:23	53261	error.html.html
2026-09-03 15:35:58	53261	error.html.txt
2026-09-03 17:50:05	1073741824	large-file-1GB.bin

Clean up the 1 GB

`aws s3 rm s3://test-1342/large-file-1GB.bin`

`aws s3 ls s3://test-1342/`

delete your local file:

`rm large-file-1GB.bin`

S3-TASK

```
HP@Ashish MINGW64 ~/s3-cli-test
$ df -h
Filesystem                Size      Used Avail Use% Mounted on
C:/Program Files/Git     282G    122G   161G   44% /
D:                        165G     26G   140G   16% /d

HP@Ashish MINGW64 ~/s3-cli-test
$ dd if=/dev/zero of=large-file-1GB.bin bs=1M count=1024
1024+0 records in
1024+0 records out
1073741824 bytes (1.1 GB, 1.0 GiB) copied, 2.65017 s, 405 MB/s


HP@Ashish MINGW64 ~/s3-cli-test
$ ls -lh large-file-1GB.bin
-rw-r--r-- 1 HP 197121 1.0G Sep  3 17:49 large-file-1GB.bin

HP@Ashish MINGW64 ~/s3-cli-test
$ du -h large-file-1GB.bin
1.0G    large-file-1GB.bin

HP@Ashish MINGW64 ~/s3-cli-test
$ aws s3 ls s3://test-1342/
2026-09-03 15:56:03      162589 1aas.png
2026-09-03 17:30:18         37 cli-test.txt
2026-09-03 15:37:23      53261 error.html.html
2026-09-03 15:35:58      53261 error.html.txt
2026-09-03 17:29:46        19 hello.txt
2026-09-03 14:49:33    5378539 windows-95-dark-3840x2160-27097.jpg

HP@Ashish MINGW64 ~/s3-cli-test
$ aws s3 cp large-file-1GB.bin s3://test-1342/
upload: ./large-file-1GB.bin to s3://test-1342/large-file-1GB.bin

HP@Ashish MINGW64 ~/s3-cli-test
$ |
```



test-1342 [Info](#)

[Objects](#) [Metadata](#) [Properties](#) [Permissions](#) [Metrics](#) [Management](#) [File systems - new](#) [Access Points](#)

Objects (1)

Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 inventory](#) to get a list of all objects in your bucket. For others to access your objects, you'll need to explicitly grant them permissions. [Learn more](#)

☐ Show versions

☐	Name	Type	Last modified	Size	Storage class
☐	large-file-1GB.bin	bin	September 3, 2026, 17:50:05 (UTC+05:30)	1.0 GB	Standard